Kubernetes

Vertical Pod Autoscaler Operator

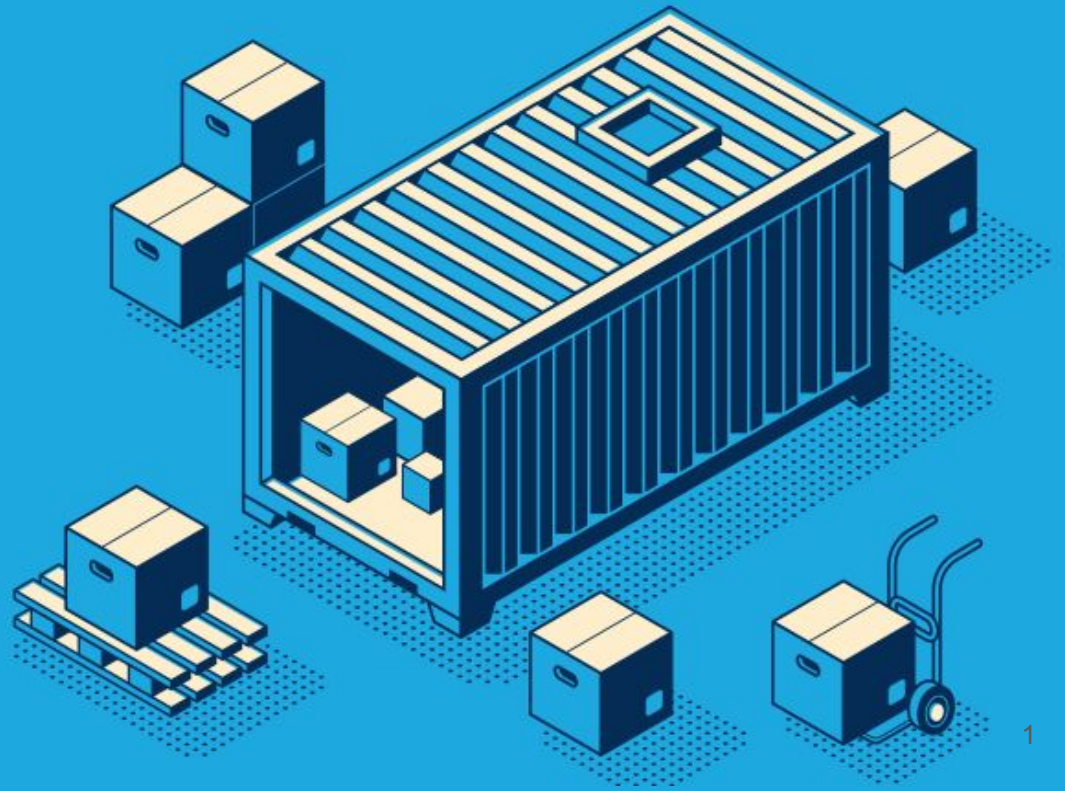

The Team

Selen

Laxmi

Apoorva

Akshay

Sindhu

Mentors (Redhat)

Humair Khan & Anand Sanmukhani

Kubernetes

An Orchestration tool for automating deployment, scaling and management of containerized applications.

Operators

Operators are software **extensions to Kubernetes** that make use of custom resources to manage applications and their components.

Without Operator

kind: Deployment
Image: frontend/v1
replicas: 1
...
spec: (Pod)

kind: Deployment
image: backend/v1
replicas: 1
...
spec: (Pod)

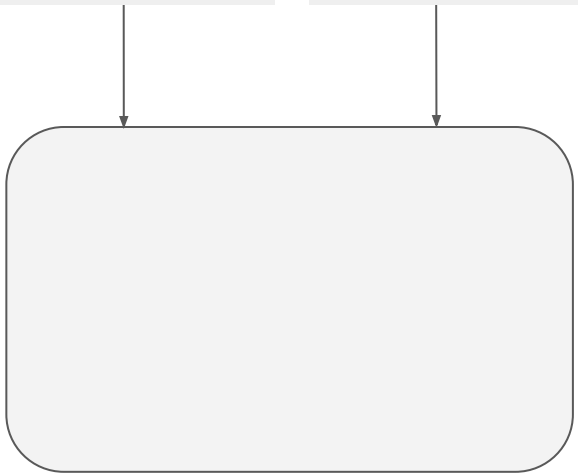

Cluster

With Operator

Without Operator

kind: Deployment
Image: frontend/v1
replicas: 1
...
spec: (Pod)

kind: Deployment
image: backend/v1
replicas: 1
...
spec: (Pod)

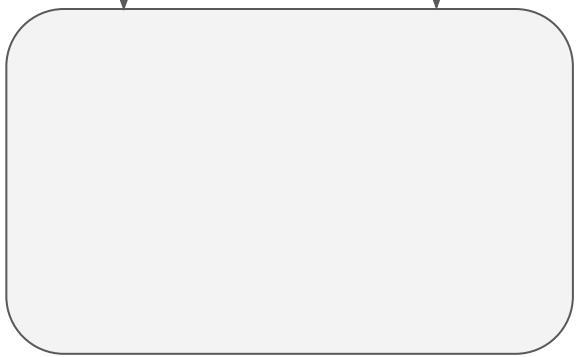

Cluster

With Operator

Controller

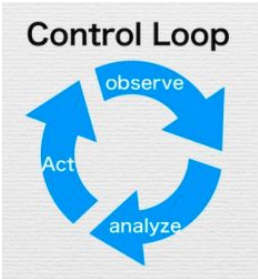

Without Operator

kind: Deployment
Image: frontend/v1
replicas: 1
...
spec: (Pod)

kind: Deployment
image: backend/v1
replicas: 1
...
spec: (Pod)

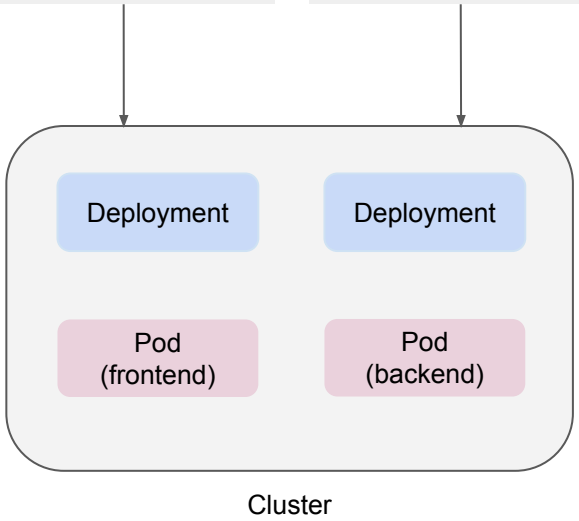

With Operator

Without Operator

kind: Deployment
Image: frontend/v1
replicas: 1
...
spec: (Pod)

kind: Deployment
image: backend/v1
replicas: 1
...
spec: (Pod)

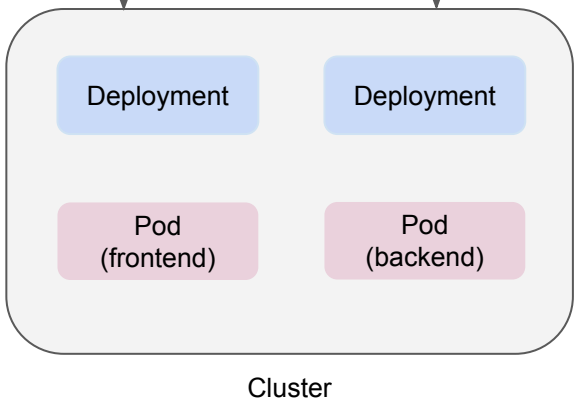

Native and
Custom
Resources

With Operator

kind: MyWebOperator
replicas: ...
...

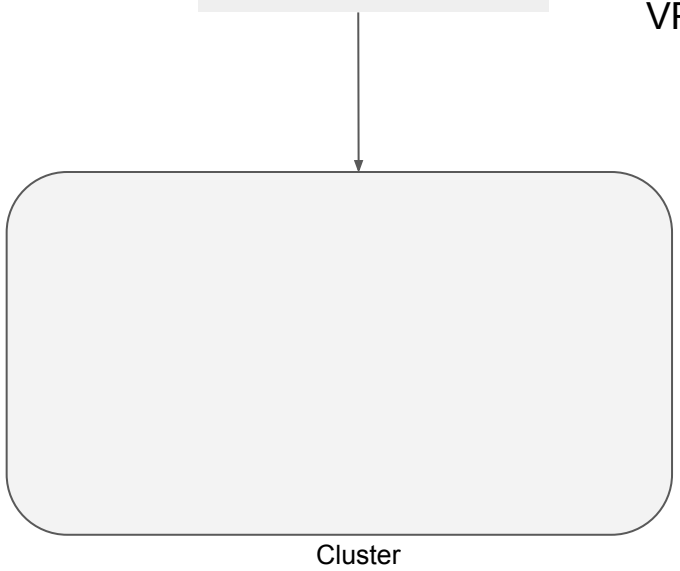

VPA

Without Operator

kind: Deployment
Image: frontend/v1
replicas: 1
...
spec: (Pod)

kind: Deployment
image: backend/v1
replicas: 1
...
spec: (Pod)

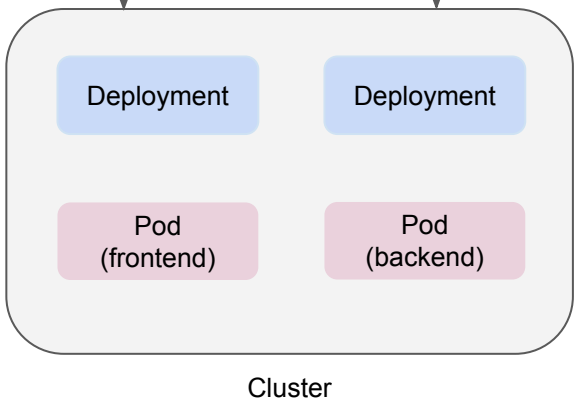

Native and
Custom
Resources

With Operator

kind: MyWebOperator
replicas: ...
...

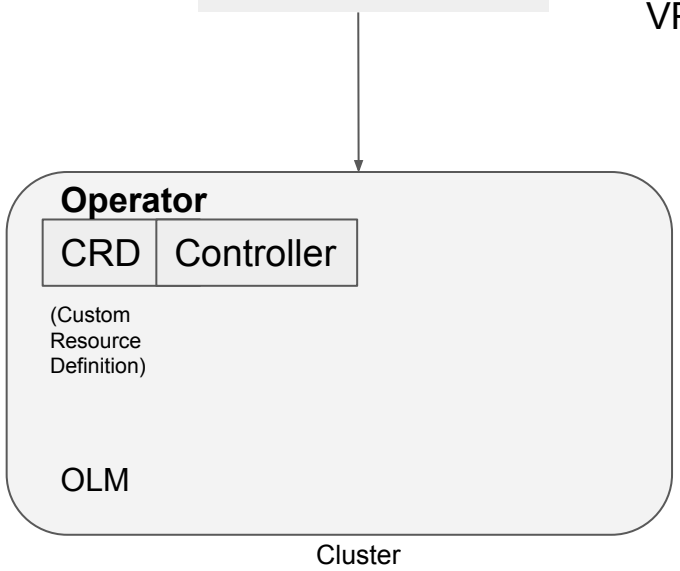

MySQL

VPA

Without Operator

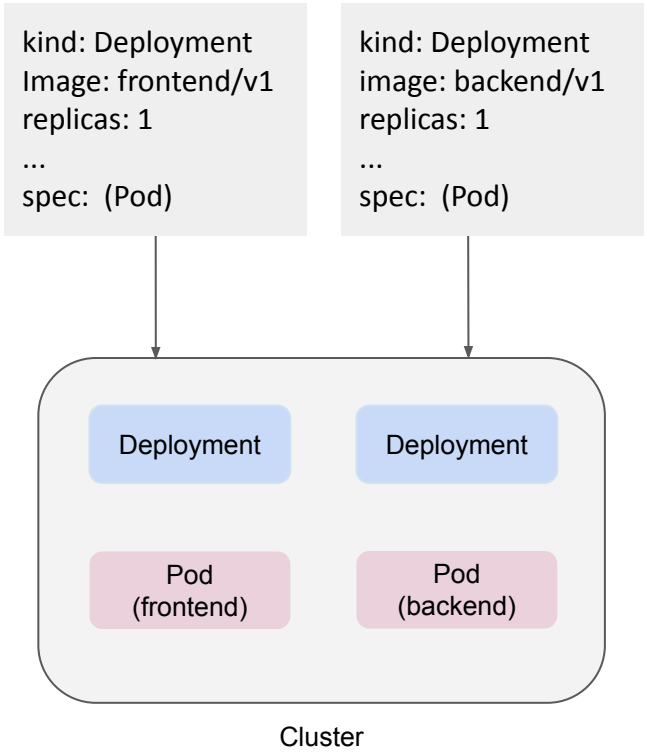

With Operator

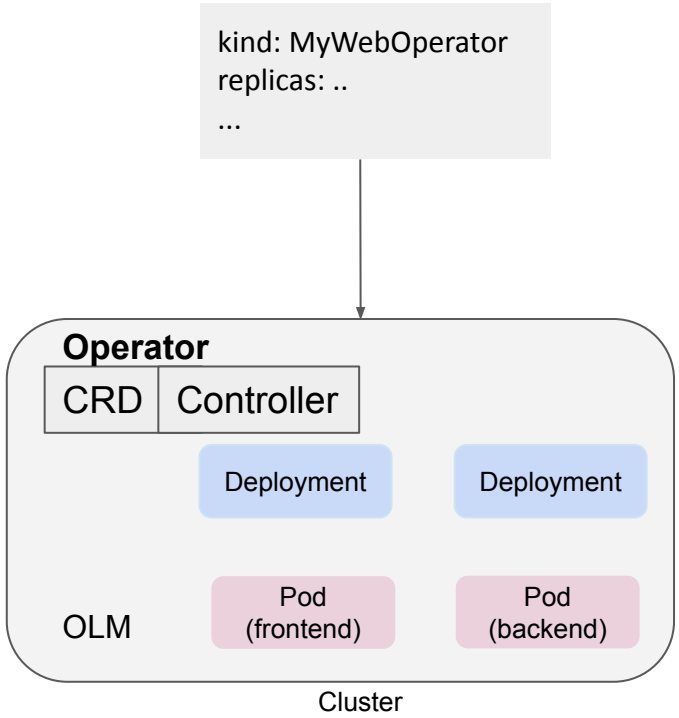

Scaling

- Horizontal Pod Autoscaling (HPA)
- Cluster Autoscaling
- Vertical Pod Autoscaling (VPA)

Container spec

```
resources:
  requests:
    memory: 1 GB
    cpu: 1 Core
  limits:
    memory: 1 GB
    cpu: 1 Core
```

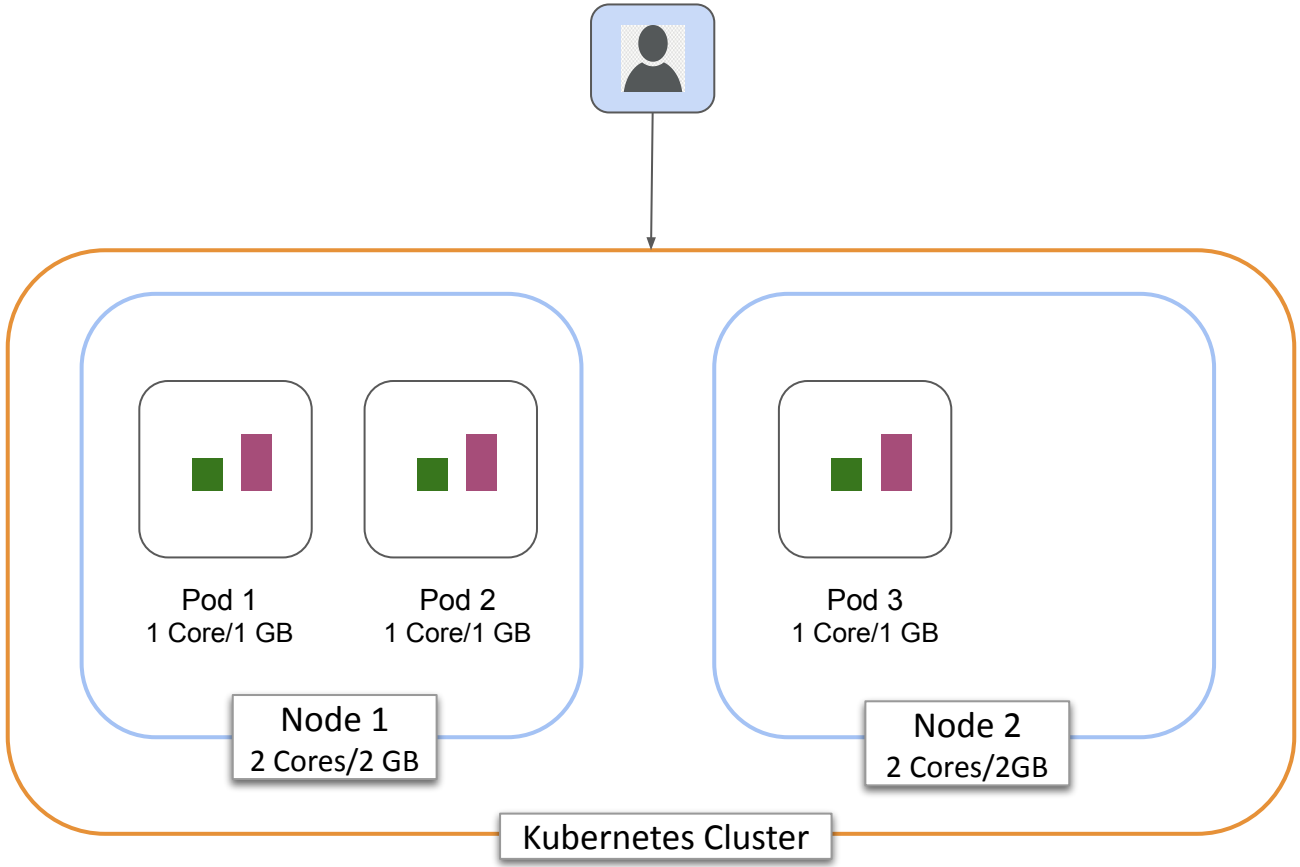

Container spec

```
resources:
  requests:
    memory: 1 GB
    cpu: 1 Core
  limits:
    memory: 1 GB
    cpu: 1 Core
```

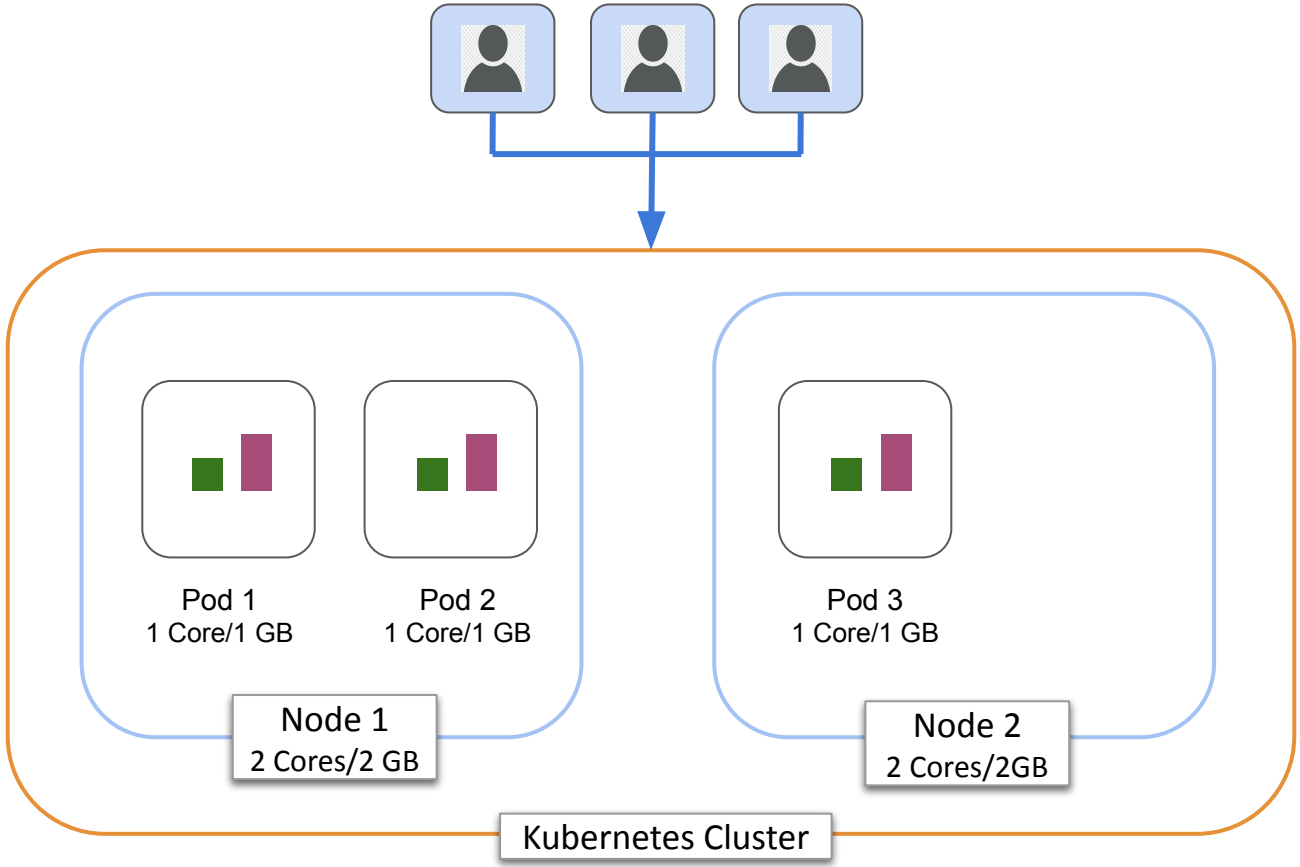

Horizontal Pod Autoscaling (HPA)

- great for handing more requests

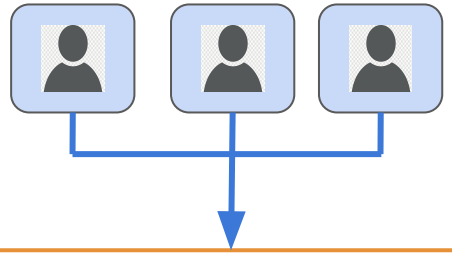

Container spec

```
resources:
  requests:
    memory: 1 GB
    cpu: 1 Core
  limits:
    memory: 1 GB
    cpu: 1 Core
```

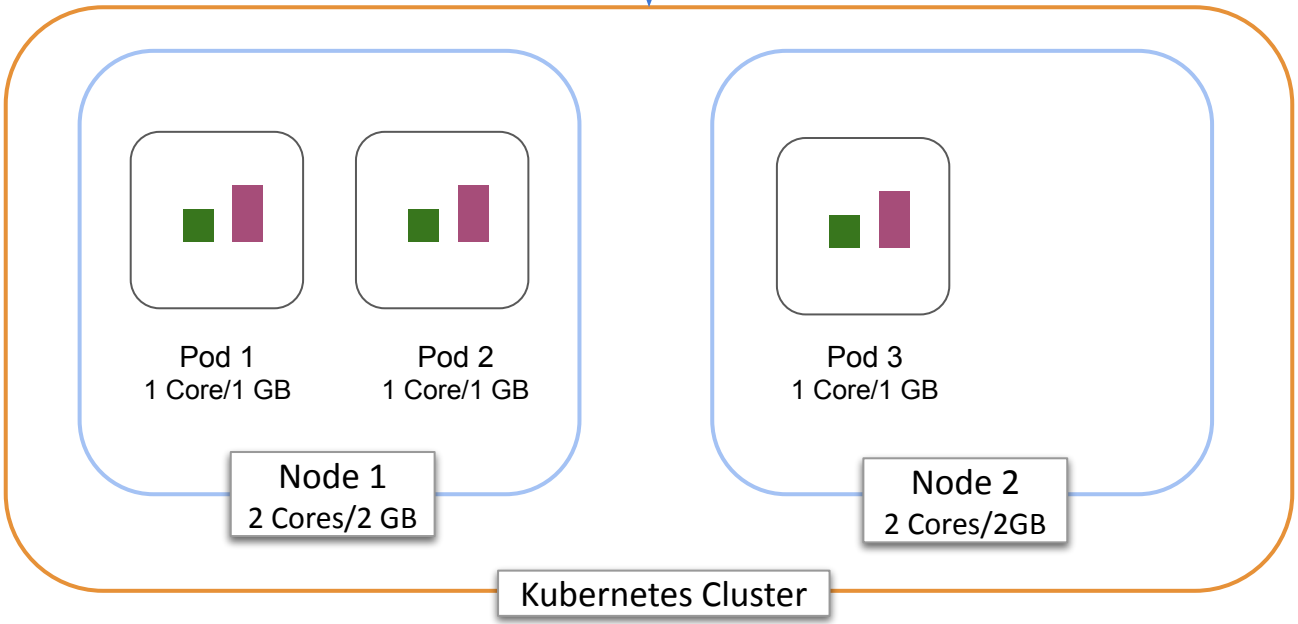

Horizontal Pod Autoscaling (HPA)

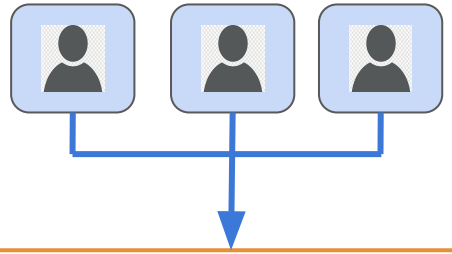

Container spec

```
resources:
  requests:
    memory: 1 GB
    cpu: 1 Core
  limits:
    memory: 1 GB
    cpu: 1 Core
```

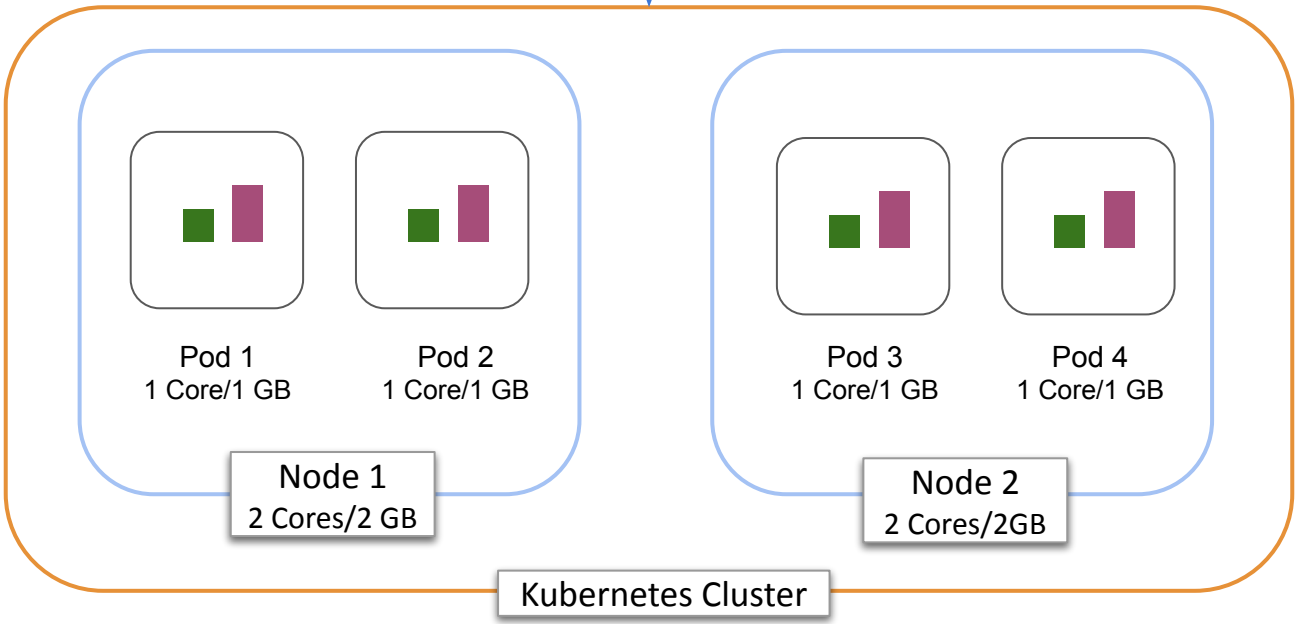

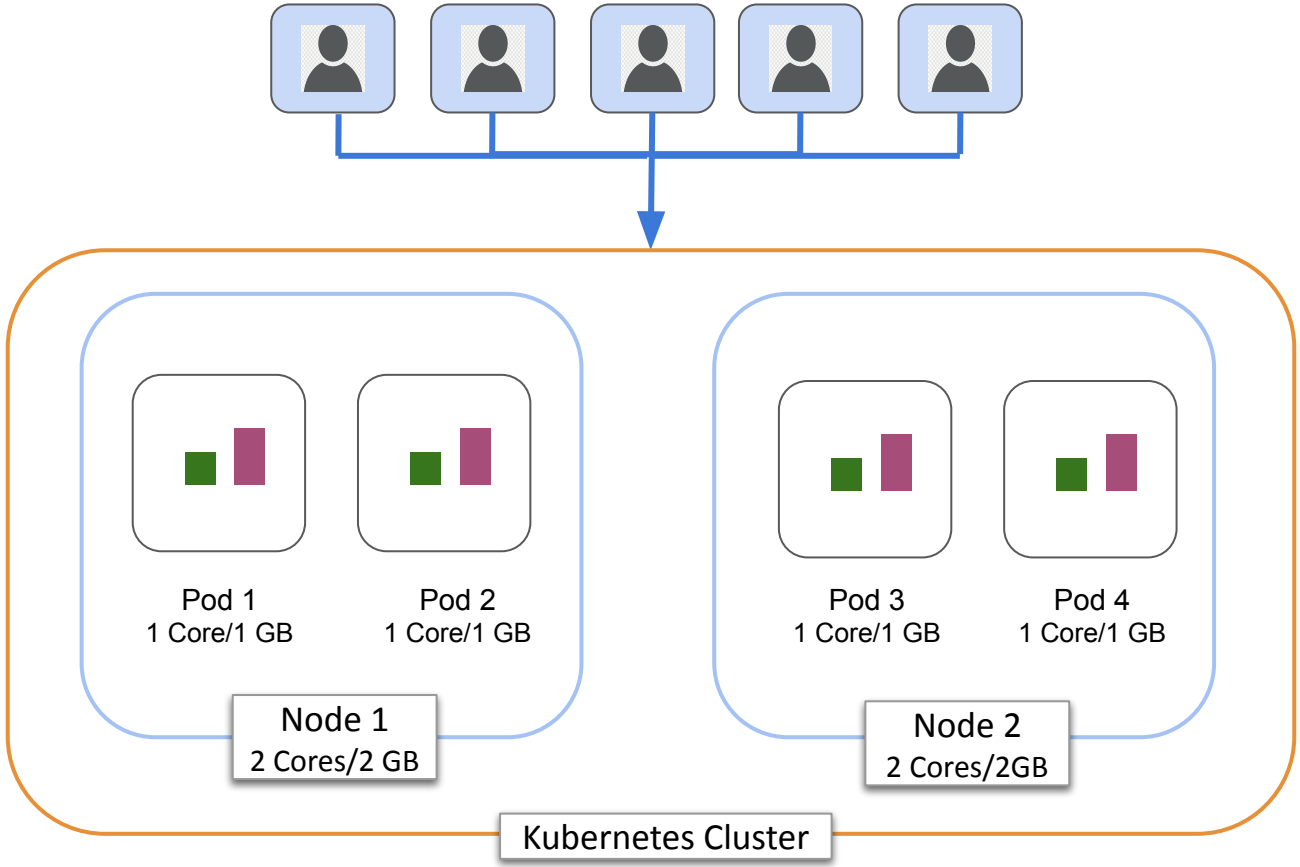
Pod 1
1 Core/1 GB
Pod 2
1 Core/1 GB
Pod 3
1 Core/1 GB
Pod 4
1 Core/1 GB
Node 1
2 Cores/2 GB
Node 2
2 Cores/2GB
Kubernetes Cluster

Cluster Autoscaling

- Scales overall cluster size

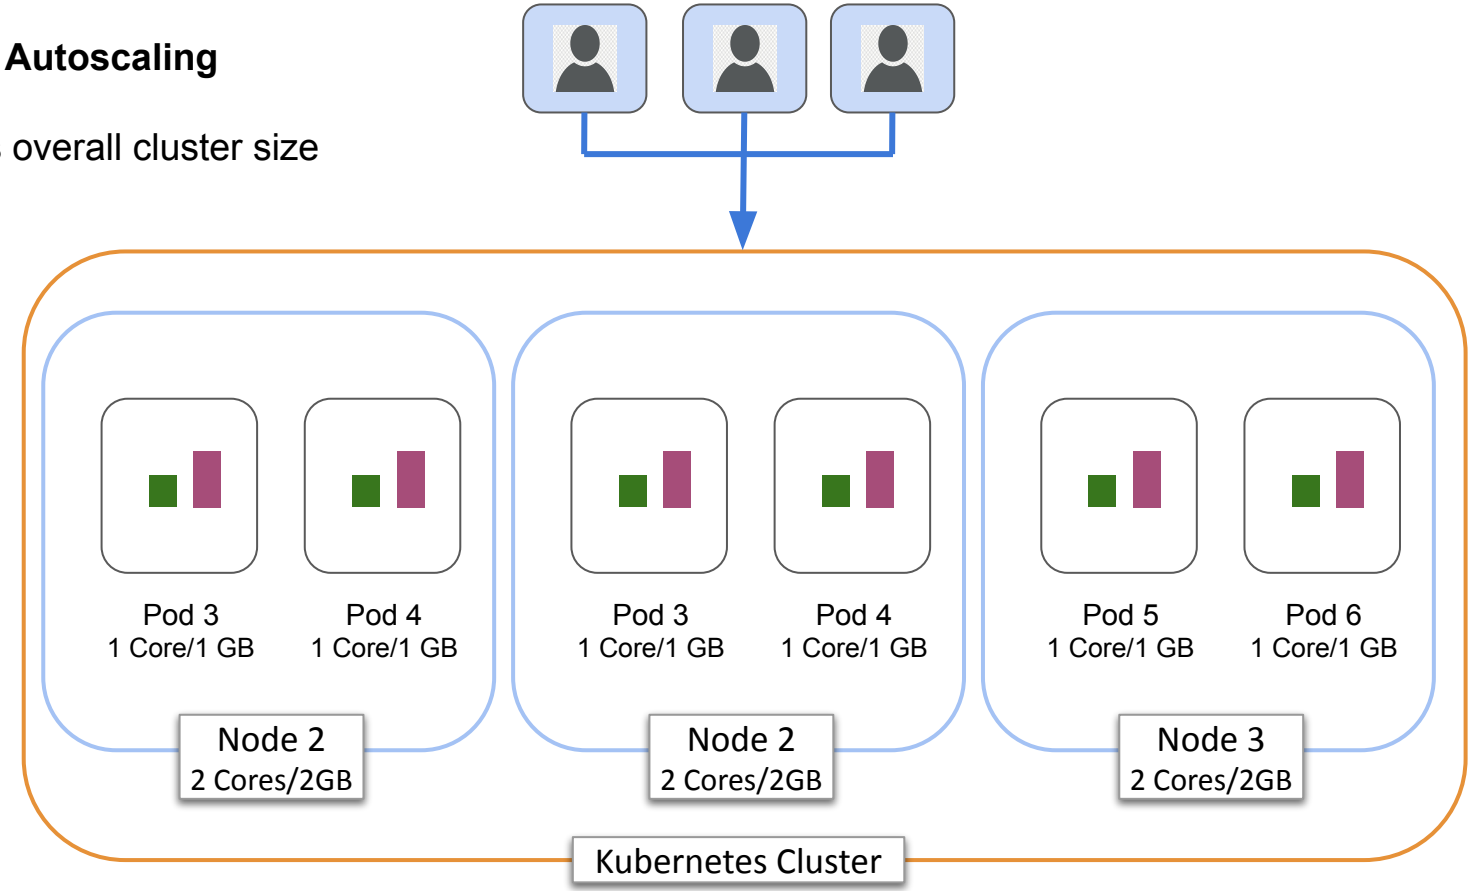

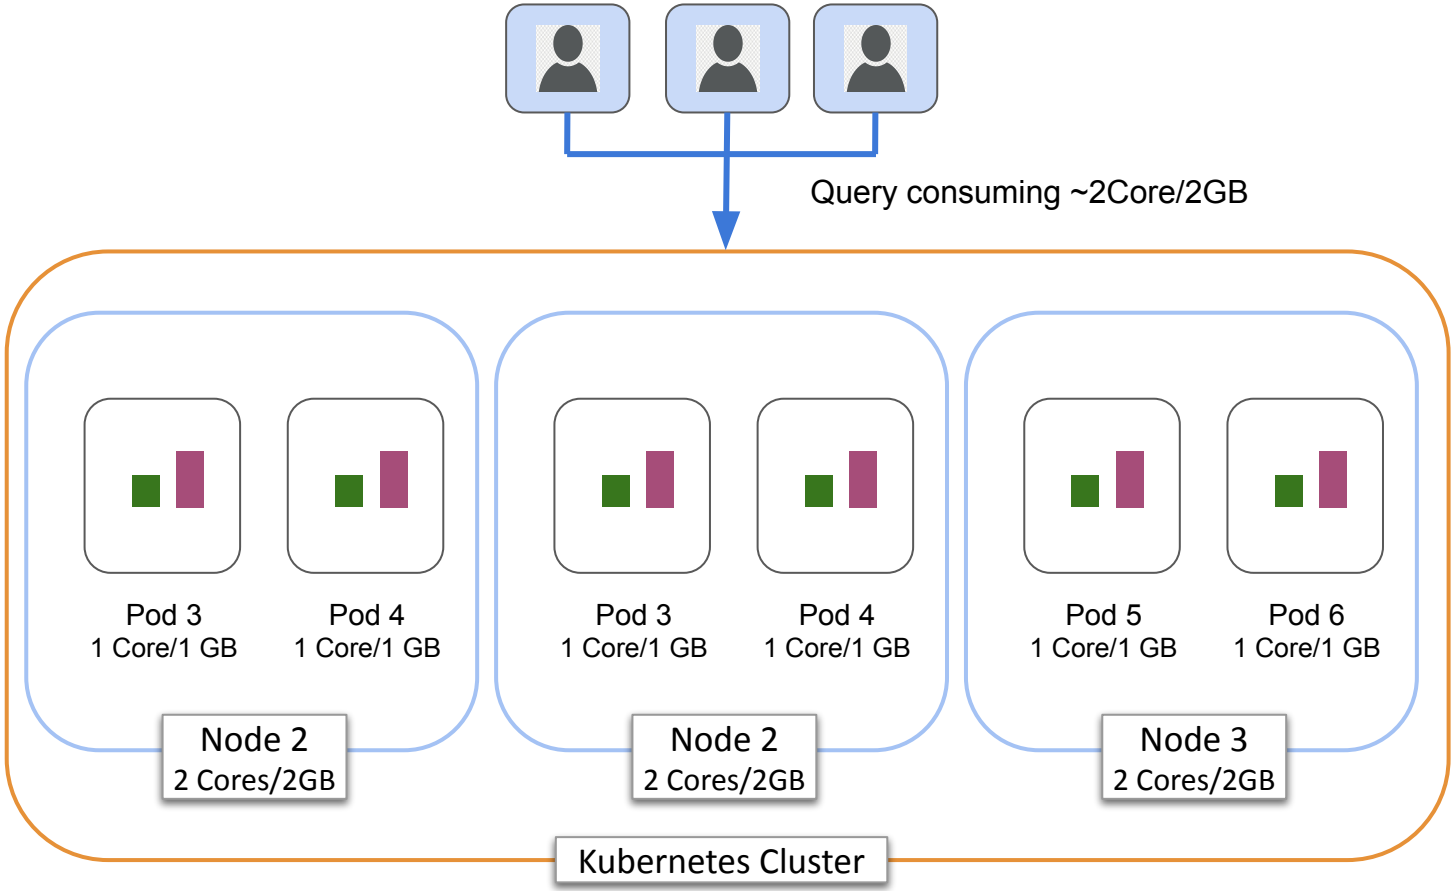
Query consuming ~2Core/2GB
Pod 3
1 Core/1 GB
Pod 4
1 Core/1 GB
Node 2
2 Cores/2GB
Pod 3
1 Core/1 GB
Pod 4
1 Core/1 GB
Node 2
2 Cores/2GB
Pod 5
1 Core/1 GB
Pod 6
1 Core/1 GB
Node 3
2 Cores/2GB
Kubernetes Cluster

Vertical Pod Autoscaling (VPA)

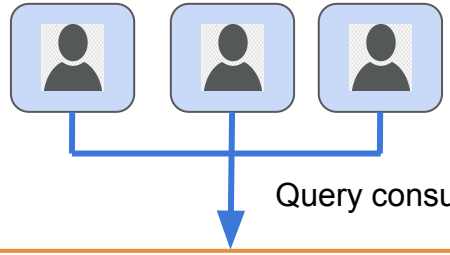

Query consuming ~2Core/2GB

VPA Recommendation

```
resources:
  requests:
    memory: 2 GB
    cpu: 2 Core
  limits:
    memory: 2 GB
    cpu: 2 Core
```

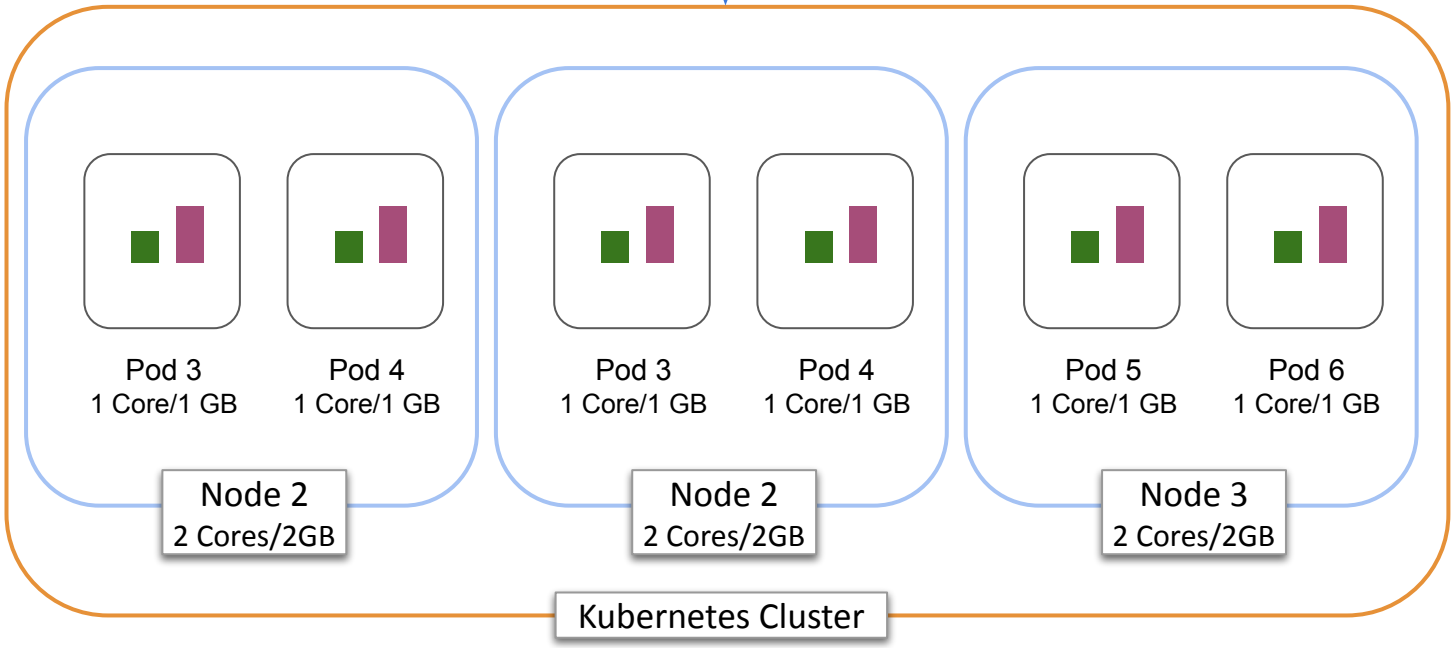

Vertical Pod Autoscaling (VPA)

VPA Recommendation

```
resources:
  requests:
    memory: 2 GB
    cpu: 2 Core
  limits:
    memory: 2 GB
    cpu: 2 Core
```

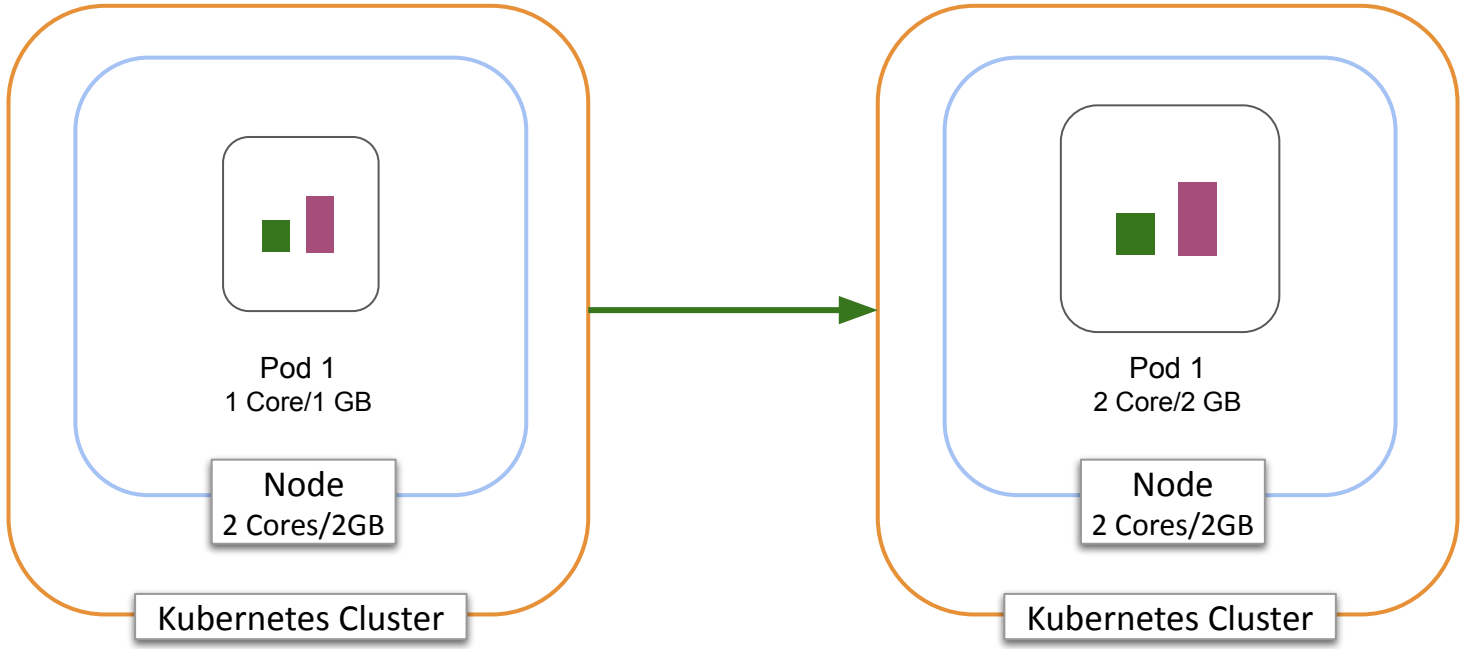

What we achieved in last sprint

(Sprint chicago)

- Installed VPA and Grafana operators in both MOC and Operate First Production cluster(Smaug).
- Deployed Resource Consumer(RC).
- Automated requests to the RC using Cron Jobs.
- Created a Grafana instance.

AUTOMATED REQUESTS TO RESOURCE CONSUMER

Manual requests via CLI

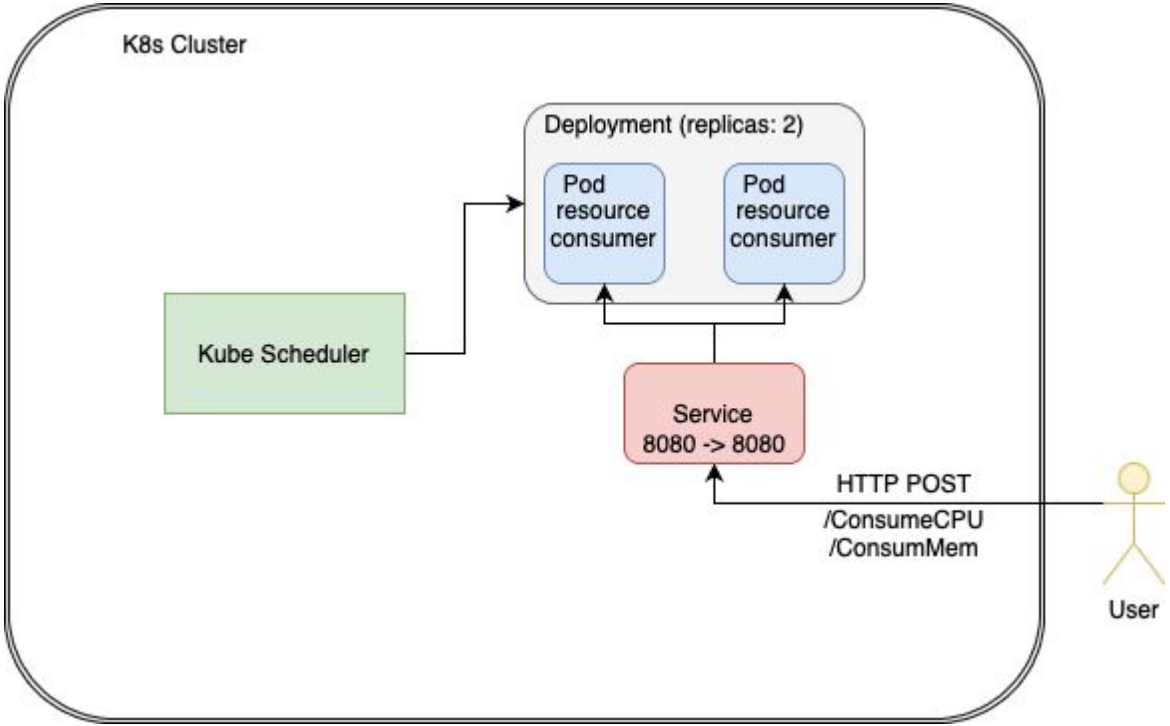

Automated requests to resource consumer

Automated requests using CronJob

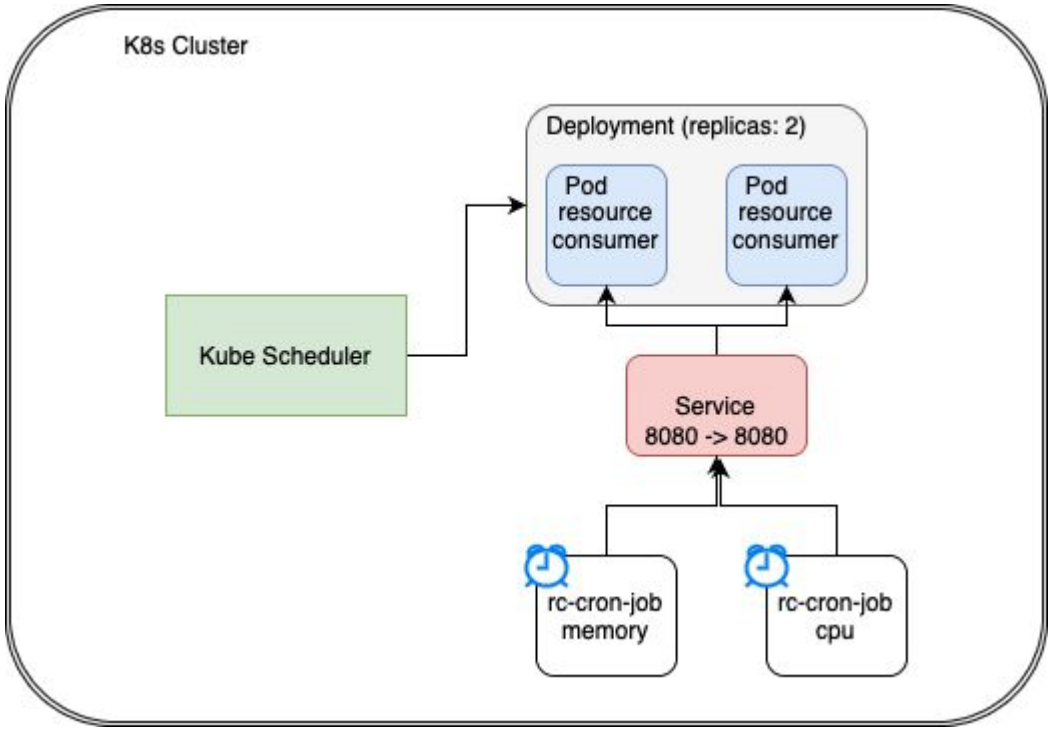

AUTOMATED REQUESTS TO RESOURCE CONSUMER

- Scheduled to send HTTP requests to /ConsumeMem and /ConsumeCPU every 10 min in the working hours (8AM - 5PM) daily.

```
spec:
  schedule: "*/10 8-17 * * *"
```

AUTOMATED REQUESTS TO RESOURCE CONSUMER

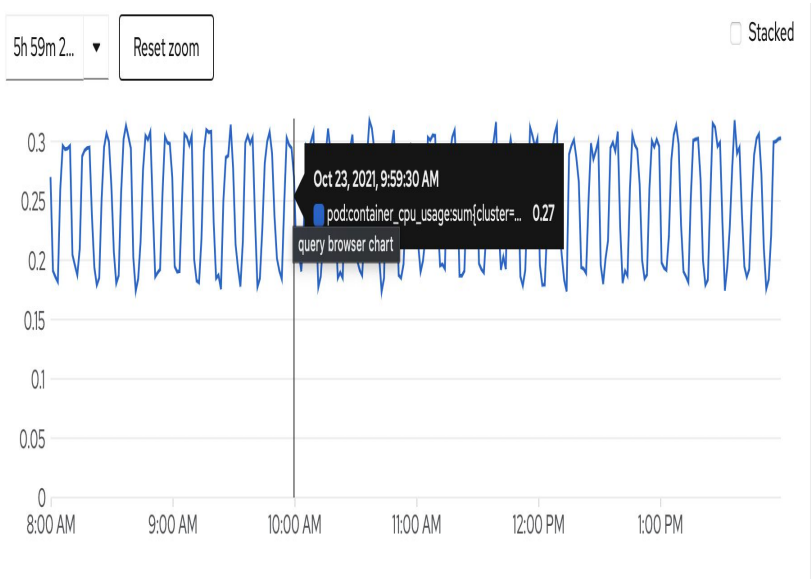

Requests from CronJob to consume CPU

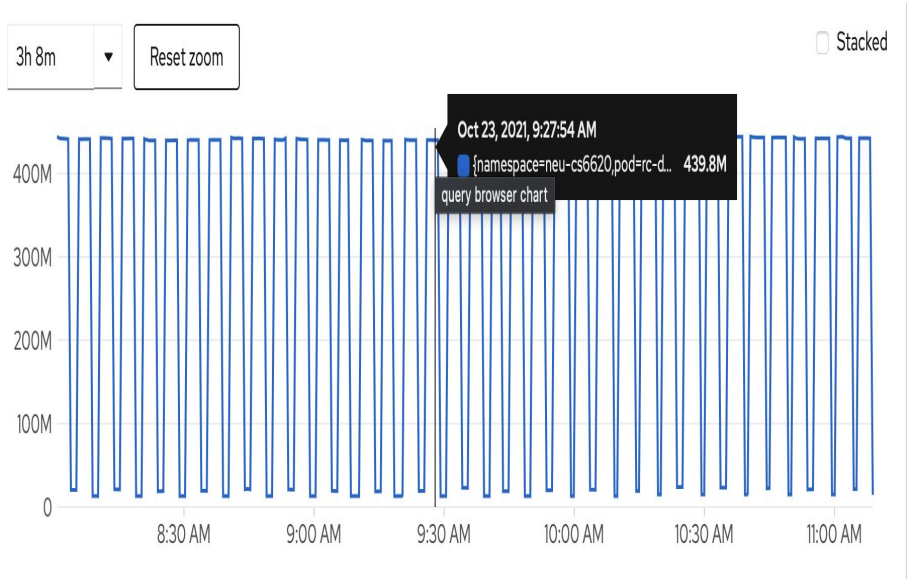

Requests from CronJob to consume Memory

AUTOMATE ROUTE CONFIGURATION

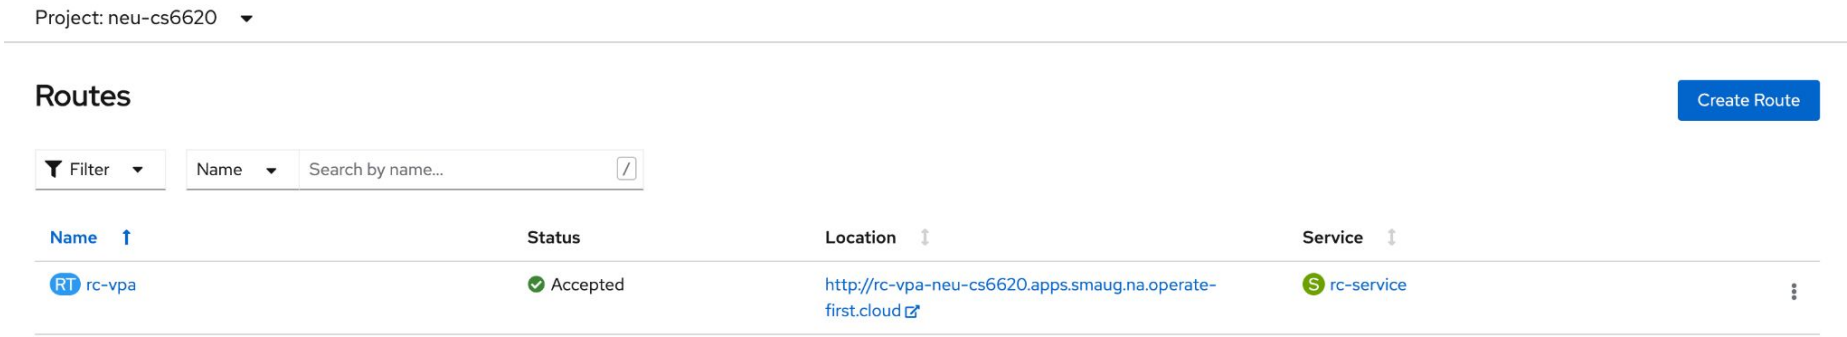

main ▾ vertical-pod-autoscaler-operator / install / rc-depl.yaml

```
...
44  kind: Route
45  apiVersion: route.openshift.io/v1
46  metadata:
47    name: rc-vpa
48  spec:
49    host: rc-vpa-neu-cs6620.apps.smaug.na.operate-first.cloud
50    to:
51      kind: Service
52      name: rc-service
53    port:
54      targetPort: 8080
55
...
```

```
$ kubectl apply -f rc-depl.yaml
```

```
$ curl --insecure --data "milicores=900&durationSec=600"
http://rc-vpa-neu-cs6620.apps.smaug.na.operate-first.cloud:8080/ConsumeCPU
```

Grafana

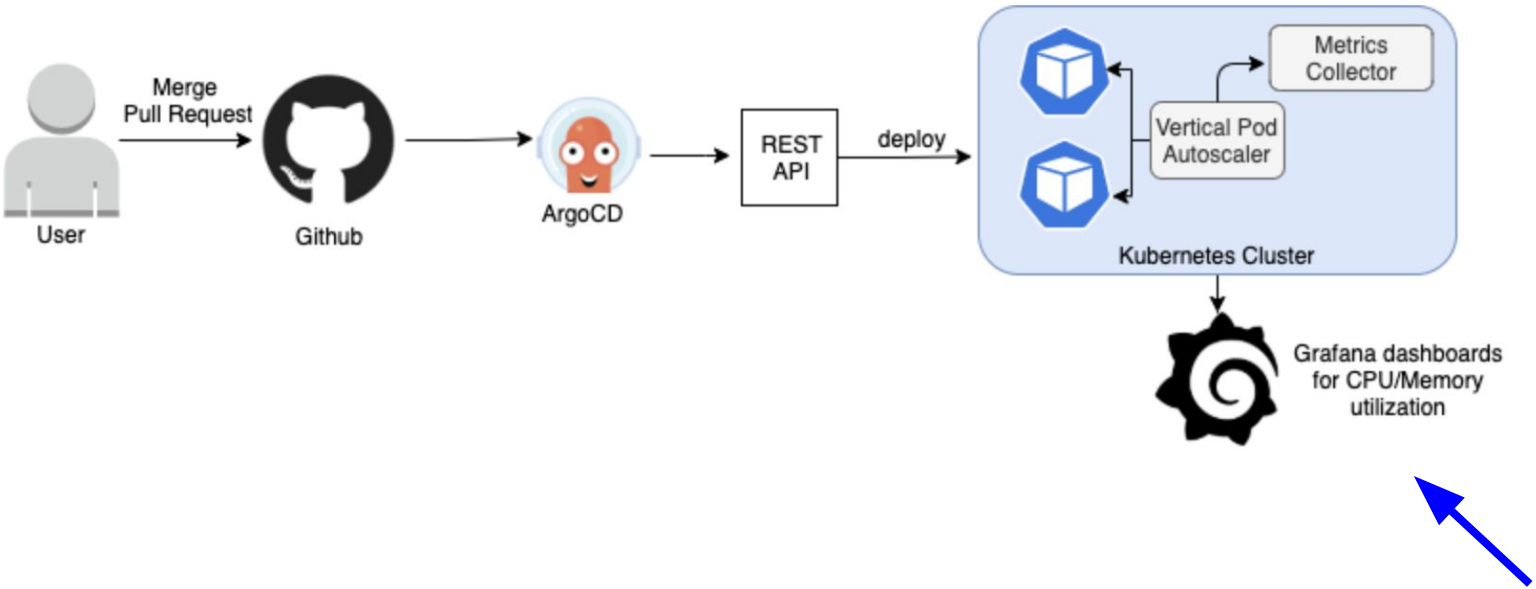
User
Merge
Pull Request
Github
ArgoCD
REST
API
deploy
Metrics
Collector
Vertical Pod
Autoscaler
Kubernetes Cluster
Grafana dashboards
for CPU/Memory
utilization

GRAFANA

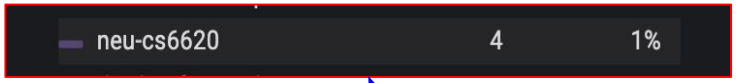
neu-cs6620
4
1%

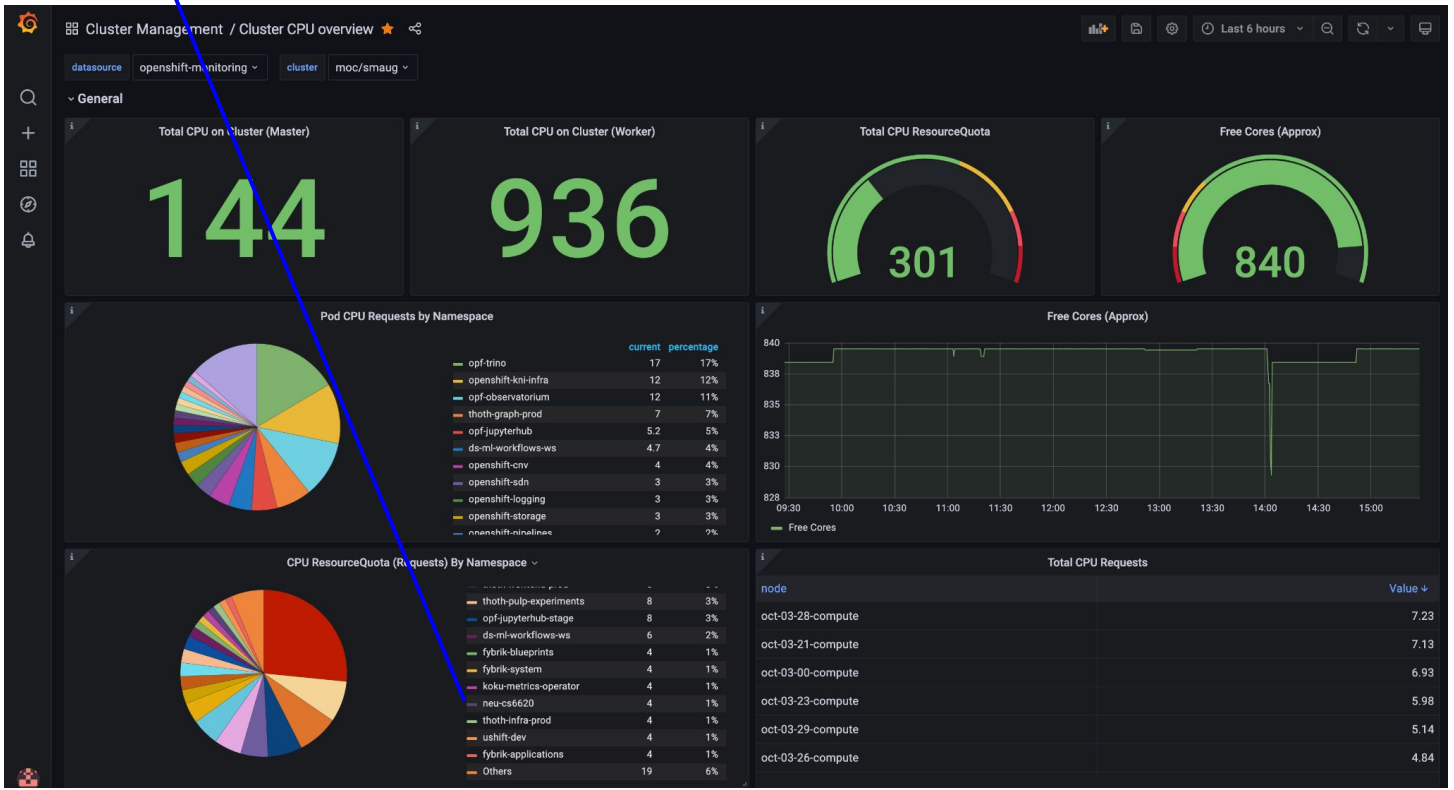
Cluster Management / Cluster CPU overview
Last 6 hours
datasource
openshift-monitoring
cluster
moc/smaug
General
Total CPU on Cluster (Master)
144
Total CPU on Cluster (Worker)
936
Total CPU ResourceQuota
301
Free Cores (Approx)
840
Pod CPU Requests by Namespace
current
percentage
opf-trino 17 17%
openshift-kni-infra 12 12%
opf-observatorium 12 11%
thoth-graph-prod 7 7%
opf-jupyterhub 5.2 5%
ds-ml-workflows-ws 4.7 4%
openshift-cnv 4 4%
openshift-sdn 3 3%
openshift-logging 3 3%
openshift-storage 3 3%
Free Cores (Approx)
840
838
835
833
830
828
09:30
10:00
10:30
11:00
11:30
12:00
12:30
13:00
13:30
14:00
14:30
15:00
Free Cores
CPU ResourceQuota (Requests) By Namespace
thoth-pulp-experiments 8 3%
opf-jupyterhub-stage 8 3%
ds-ml-workflows-ws 6 2%
fybrik-blueprints 4 1%
fybrik-system 4 1%
koku-metrics-operator 4 1%
neu-cs6620 4 1%
thoth-infra-prod 4 1%
ushift-dev 4 1%
fybrik-applications 4 1%
Others 19 6%
Total CPU Requests
node
Value
oct-03-28-compute 7.23
oct-03-21-compute 7.13
oct-03-00-compute 6.93
oct-03-23-compute 5.98
oct-03-29-compute 5.14
oct-03-26-compute 4.84

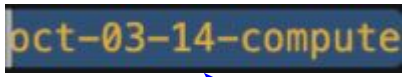

Project: neu-cs6620

Pods > Pod details

rc-replica-controller-cmlqm Running

Details Metrics YAML Environment Logs Events Terminal

```
139       name: rc-replica-controller
140       uid: 93b98848-e659-47a9-b121-b14f42ecd993
141       controller: true
142       blockOwnerDeletion: true
143   labels:
144     app: httpd
145 spec:
146   restartPolicy: Always
147   serviceAccountName: default
148   imagePullSecrets:
149     - name: default-dockercfg-zg5mp
150   priority: 0
151   schedulerName: default-scheduler
152   enableServiceLinks: true
153   terminationGracePeriodSeconds: 30
154   preemptionPolicy: PreemptLowerPriority
155   nodeName: oct-03-14-compute
156   securityContext:
157     seLinuxOptions:
158       level: 's0:c33,c32'
159     fsGroup: 1001120000
160   containers:
161     - resources:
162         limits:
163           cpu: 500m
164           memory: 1000Mi
165         requests:
166           cpu: 300m
167           memory: 400Mi
168       terminationMessagePath: /dev/termination-log
169       name: httpd
170       securityContext:
171         capabilities:
```

Save Reload Cancel

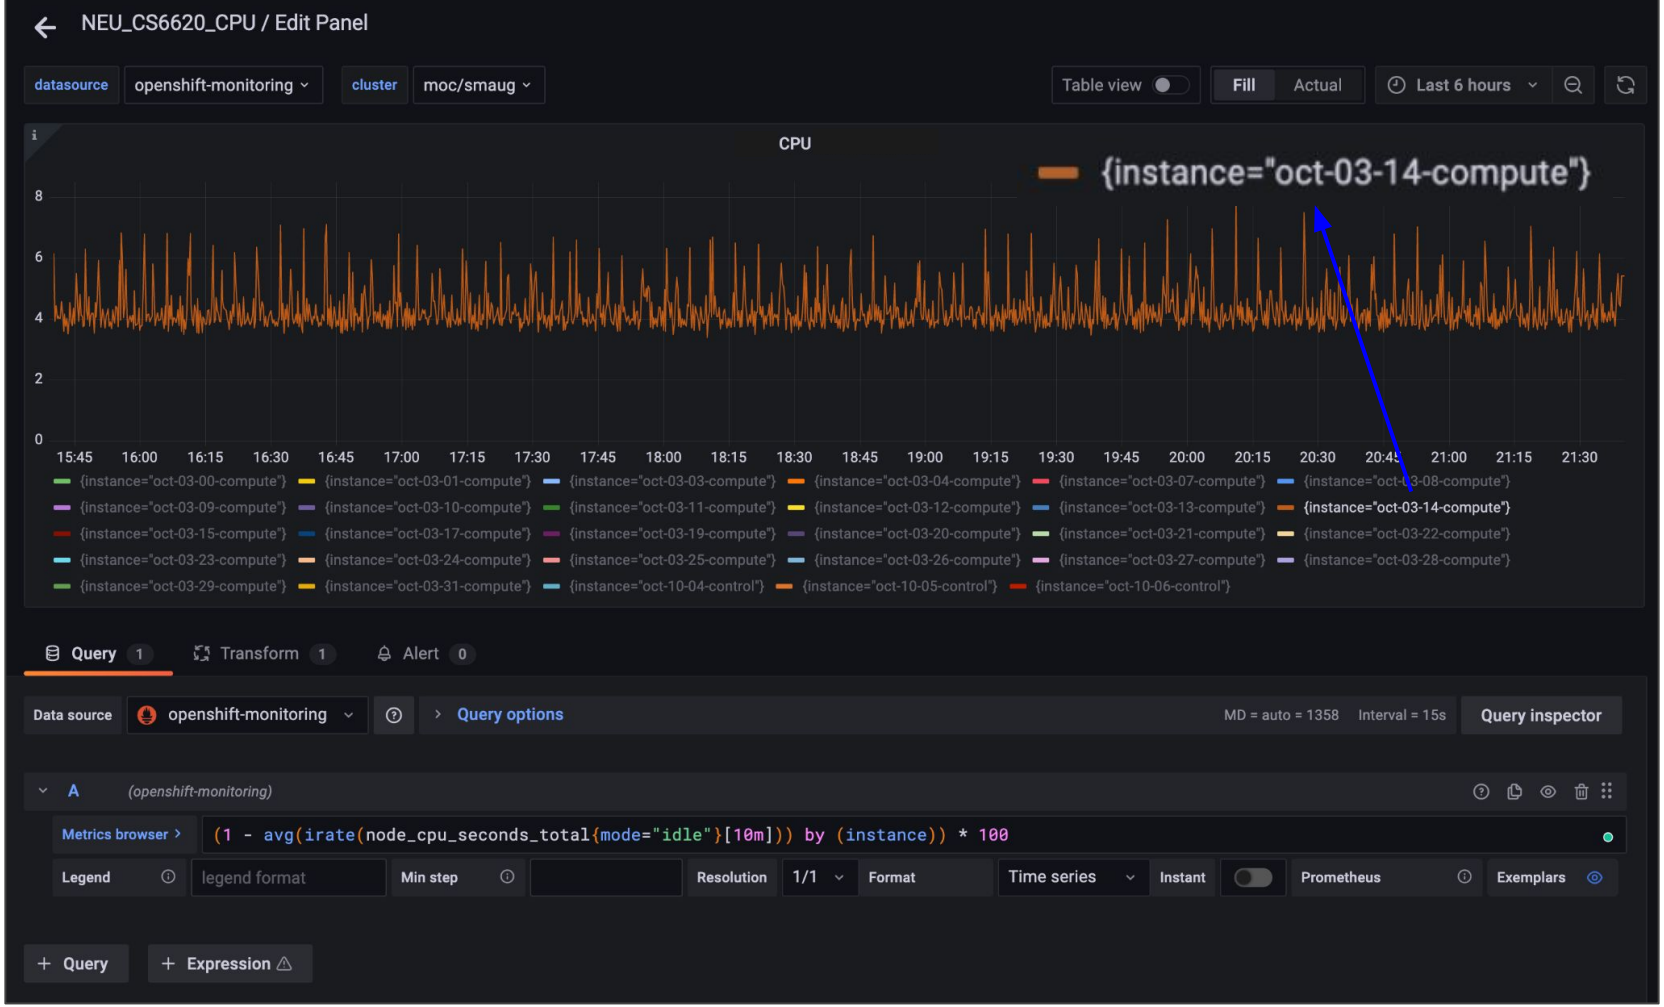
NEU_CS6620_CPU / Edit Panel
datasource
openshift-monitoring
cluster
moc/smaug
Table view
Fill
Actual
Last 6 hours
CPU
{instance="oct-03-14-compute"}
8
6
4
2
0
15:45
16:00
16:15
16:30
16:45
17:00
17:15
17:30
17:45
18:00
18:15
18:30
18:45
19:00
19:15
19:30
19:45
20:00
20:15
20:30
20:45
21:00
21:15
21:30
{instance="oct-03-00-compute"}
{instance="oct-03-01-compute"}
{instance="oct-03-03-compute"}
{instance="oct-03-04-compute"}
{instance="oct-03-07-compute"}
{instance="oct-03-08-compute"}
{instance="oct-03-09-compute"}
{instance="oct-03-10-compute"}
{instance="oct-03-11-compute"}
{instance="oct-03-12-compute"}
{instance="oct-03-13-compute"}
{instance="oct-03-14-compute"}
{instance="oct-03-15-compute"}
{instance="oct-03-17-compute"}
{instance="oct-03-19-compute"}
{instance="oct-03-20-compute"}
{instance="oct-03-21-compute"}
{instance="oct-03-22-compute"}
{instance="oct-03-23-compute"}
{instance="oct-03-24-compute"}
{instance="oct-03-25-compute"}
{instance="oct-03-26-compute"}
{instance="oct-03-27-compute"}
{instance="oct-03-28-compute"}
{instance="oct-03-29-compute"}
{instance="oct-03-31-compute"}
{instance="oct-10-04-control"}
{instance="oct-10-05-control"}
{instance="oct-10-06-control"}
Query 1
Transform 1
Alert 0
Data source
openshift-monitoring
Query options
MD = auto = 1358
Interval = 15s
Query inspector
A
(openshift-monitoring)
Metrics browser
(1 - avg(irate(node_cpu_seconds_total{mode="idle"}[10m])) by (instance)) * 100
Legend
legend format
Min step
Resolution
1/1
Format
Time series
Instant
Prometheus
Exemplars
+ Query
+ Expression

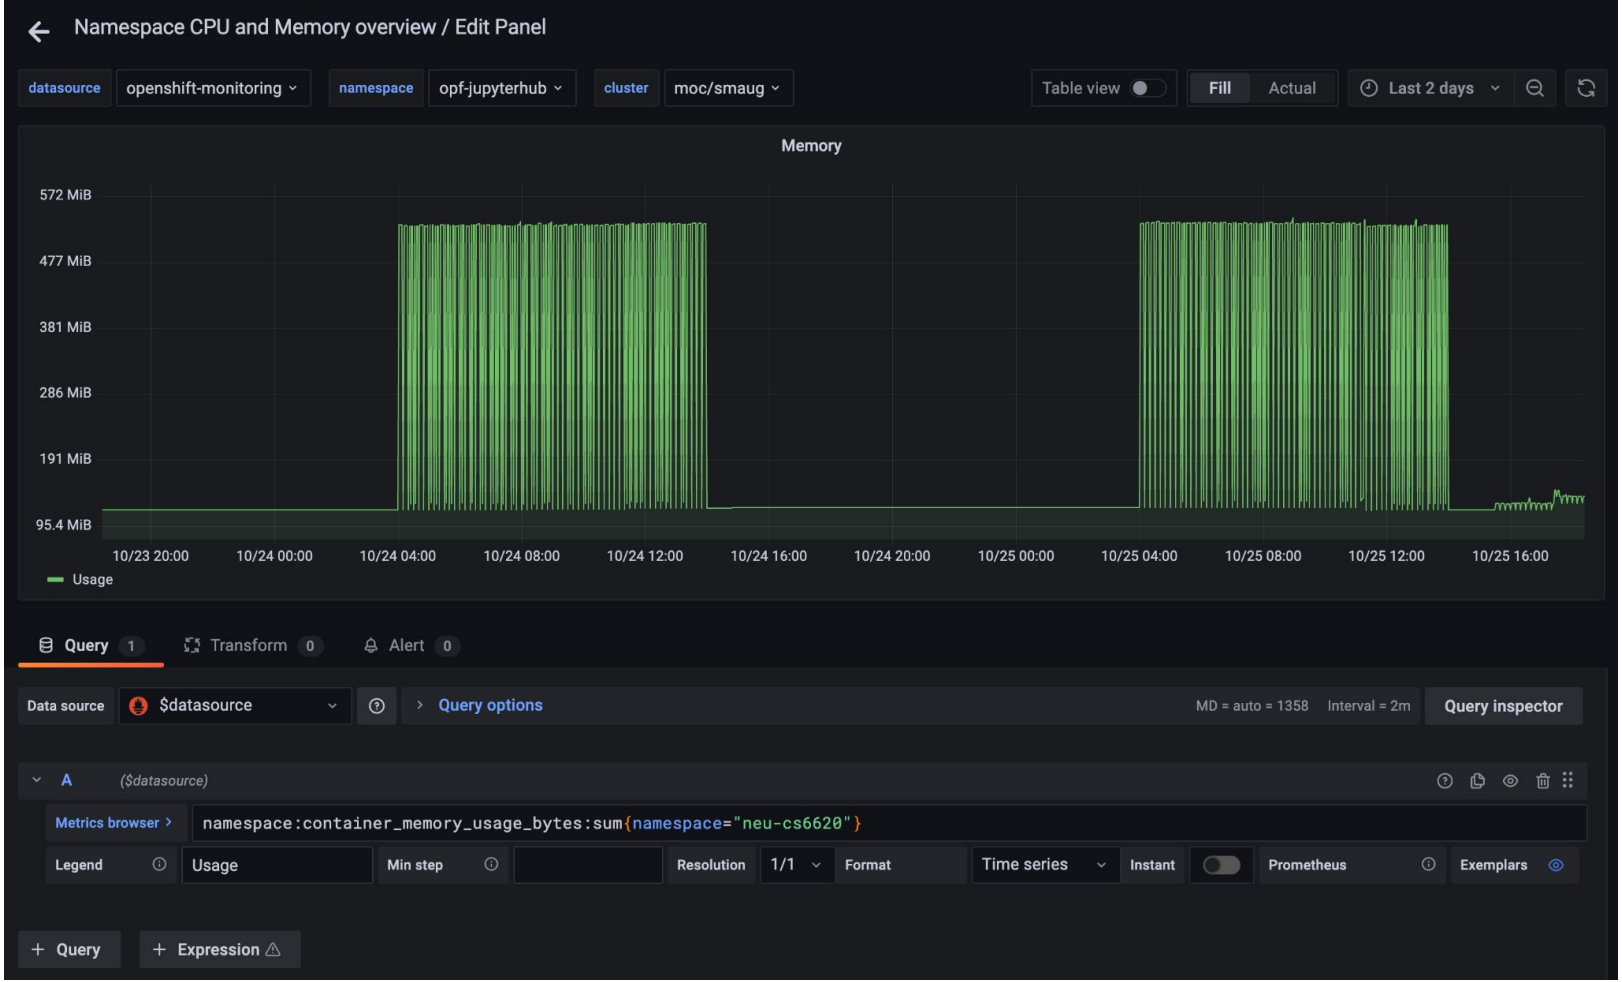
Namespace CPU and Memory overview / Edit Panel
datasource
openshift-monitoring
namespace
opf-jupyterhub
cluster
moc/smaug
Table view
Fill
Actual
Last 2 days
Memory
572 MiB
477 MiB
381 MiB
286 MiB
191 MiB
95.4 MiB
10/23 20:00
10/24 00:00
10/24 04:00
10/24 08:00
10/24 12:00
10/24 16:00
10/24 20:00
10/25 00:00
10/25 04:00
10/25 08:00
10/25 12:00
10/25 16:00
Usage
Query 1
Transform 0
Alert 0
Data source
$datasource
Query options
MD = auto = 1358
Interval = 2m
Query inspector
A
($datasource)
Metrics browser
namespace:container_memory_usage_bytes:sum{namespace="neu-cs6620"}
Legend
Usage
Min step
Resolution
1/1
Format
Time series
Instant
Prometheus
Exemplars
Query
Expression

VPA Recommendation

VPA Custom Resource

```yaml
apiVersion: "autoscaling.k8s.io/v1beta2"
kind: VerticalPodAutoscaler
metadata:
  name: resource-consumer
  namespace: cs6620-fall21-deployverticalpod
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: Deployment
    name: rc-deployment
  updatePolicy:
    updateMode: "Auto"
```

Recommendation:

```
Status:
  Conditions:
    Last Transition Time:  2021-10-25T20:47:20Z
    Status:                True
    Type:                  RecommendationProvided
  Recommendation:
    Container Recommendations:
      Container Name:  resource-consumer
      Lower Bound:
        Cpu:     25m
        Memory:  380082209
      Target:
        Cpu:     247m
        Memory:  511772986
      Uncapped Target:
        Cpu:     247m
        Memory:  511772986
      Upper Bound:
        Cpu:     505m
        Memory:  630425876
```

Resource metrics

```
resources:
  requests:
    memory: 1 GB
    cpu: 1 Core
  limits:
    memory: 1 GB
    cpu: 1 Core
```

Limit

The kubelet enforces those ***limits*** so that the running container is not allowed to use more of that resource than the limit you set.

Request

The kubelet also reserves at least the ***request*** amount of that system resource specifically for that container to use

VPA Recommendation Terms

- **lowerBound** - the minimum recommended resource levels.
- **target** - the recommended resource levels.
- **upperBound** - the highest recommended resource levels.
- **uncappedTarget** - the most recent resource recommendations.

```
Status:
  Conditions:
    Last Transition Time:  2021-10-25T20:47:20Z
    Status:                True
    Type:                  RecommendationProvided
  Recommendation:
    Container Recommendations:
      Container Name:  resource-consumer
      Lower Bound:
        Cpu:     25m
        Memory:  380082209
      Target:
        Cpu:     247m
        Memory:  511772986
      Uncapped Target:
        Cpu:     247m
        Memory:  511772986
      Upper Bound:
        Cpu:     505m
        Memory:  630425876
```

VPA Update policy

- **Auto** - Updates the pod based on its recommendation by deleting and recreating it.
- **Initial** - Applies recommendations only during pod creation.
- **Off** - Just analyzes and records the recommendations.

```yaml
apiVersion: "autoscaling.k8s.io/v1beta2"
kind: VerticalPodAutoscaler
metadata:
  name: resource-consumer
  namespace: cs6620-fall21-deployverticalpod
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: Deployment
    name: rc-deployment
  updatePolicy:
    updateMode: "Auto"
```

VPA Test cases

1. Pod over-utilization within upper bound of VPA recommendation.
2. Pod under-utilization within lower bound of VPA recommendations
3. VPA automatically deletes and creating a pod with its new recommendation for upperBound case (over-utilizing) within the limit
4. VPA automatically deletes and creating a pod with its new recommendation for lowerBound case (under-utilizing)

VPA Demo

Sprint Chicago 12 Oct 2021 to 25 Oct 2021

70% 80 total points 56 completed points 1 open tasks 31 closed tasks 0 iocaine doses

How this chart works

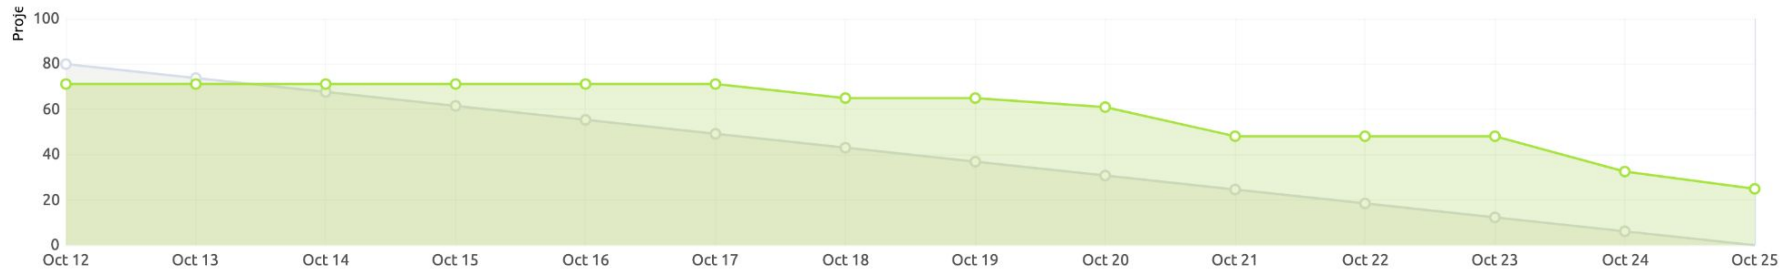

What we plan for Next sprint

(Sprint Dallas)

- Experiment VPA scenarios
  - Auto updatePolicy of VPA
  - Underutilization and Overutilization of resources
- Grafana Live Dashboard
  - Create Grafana dashboard with our metrics of interest
- Setup POC application through GitOps workflow